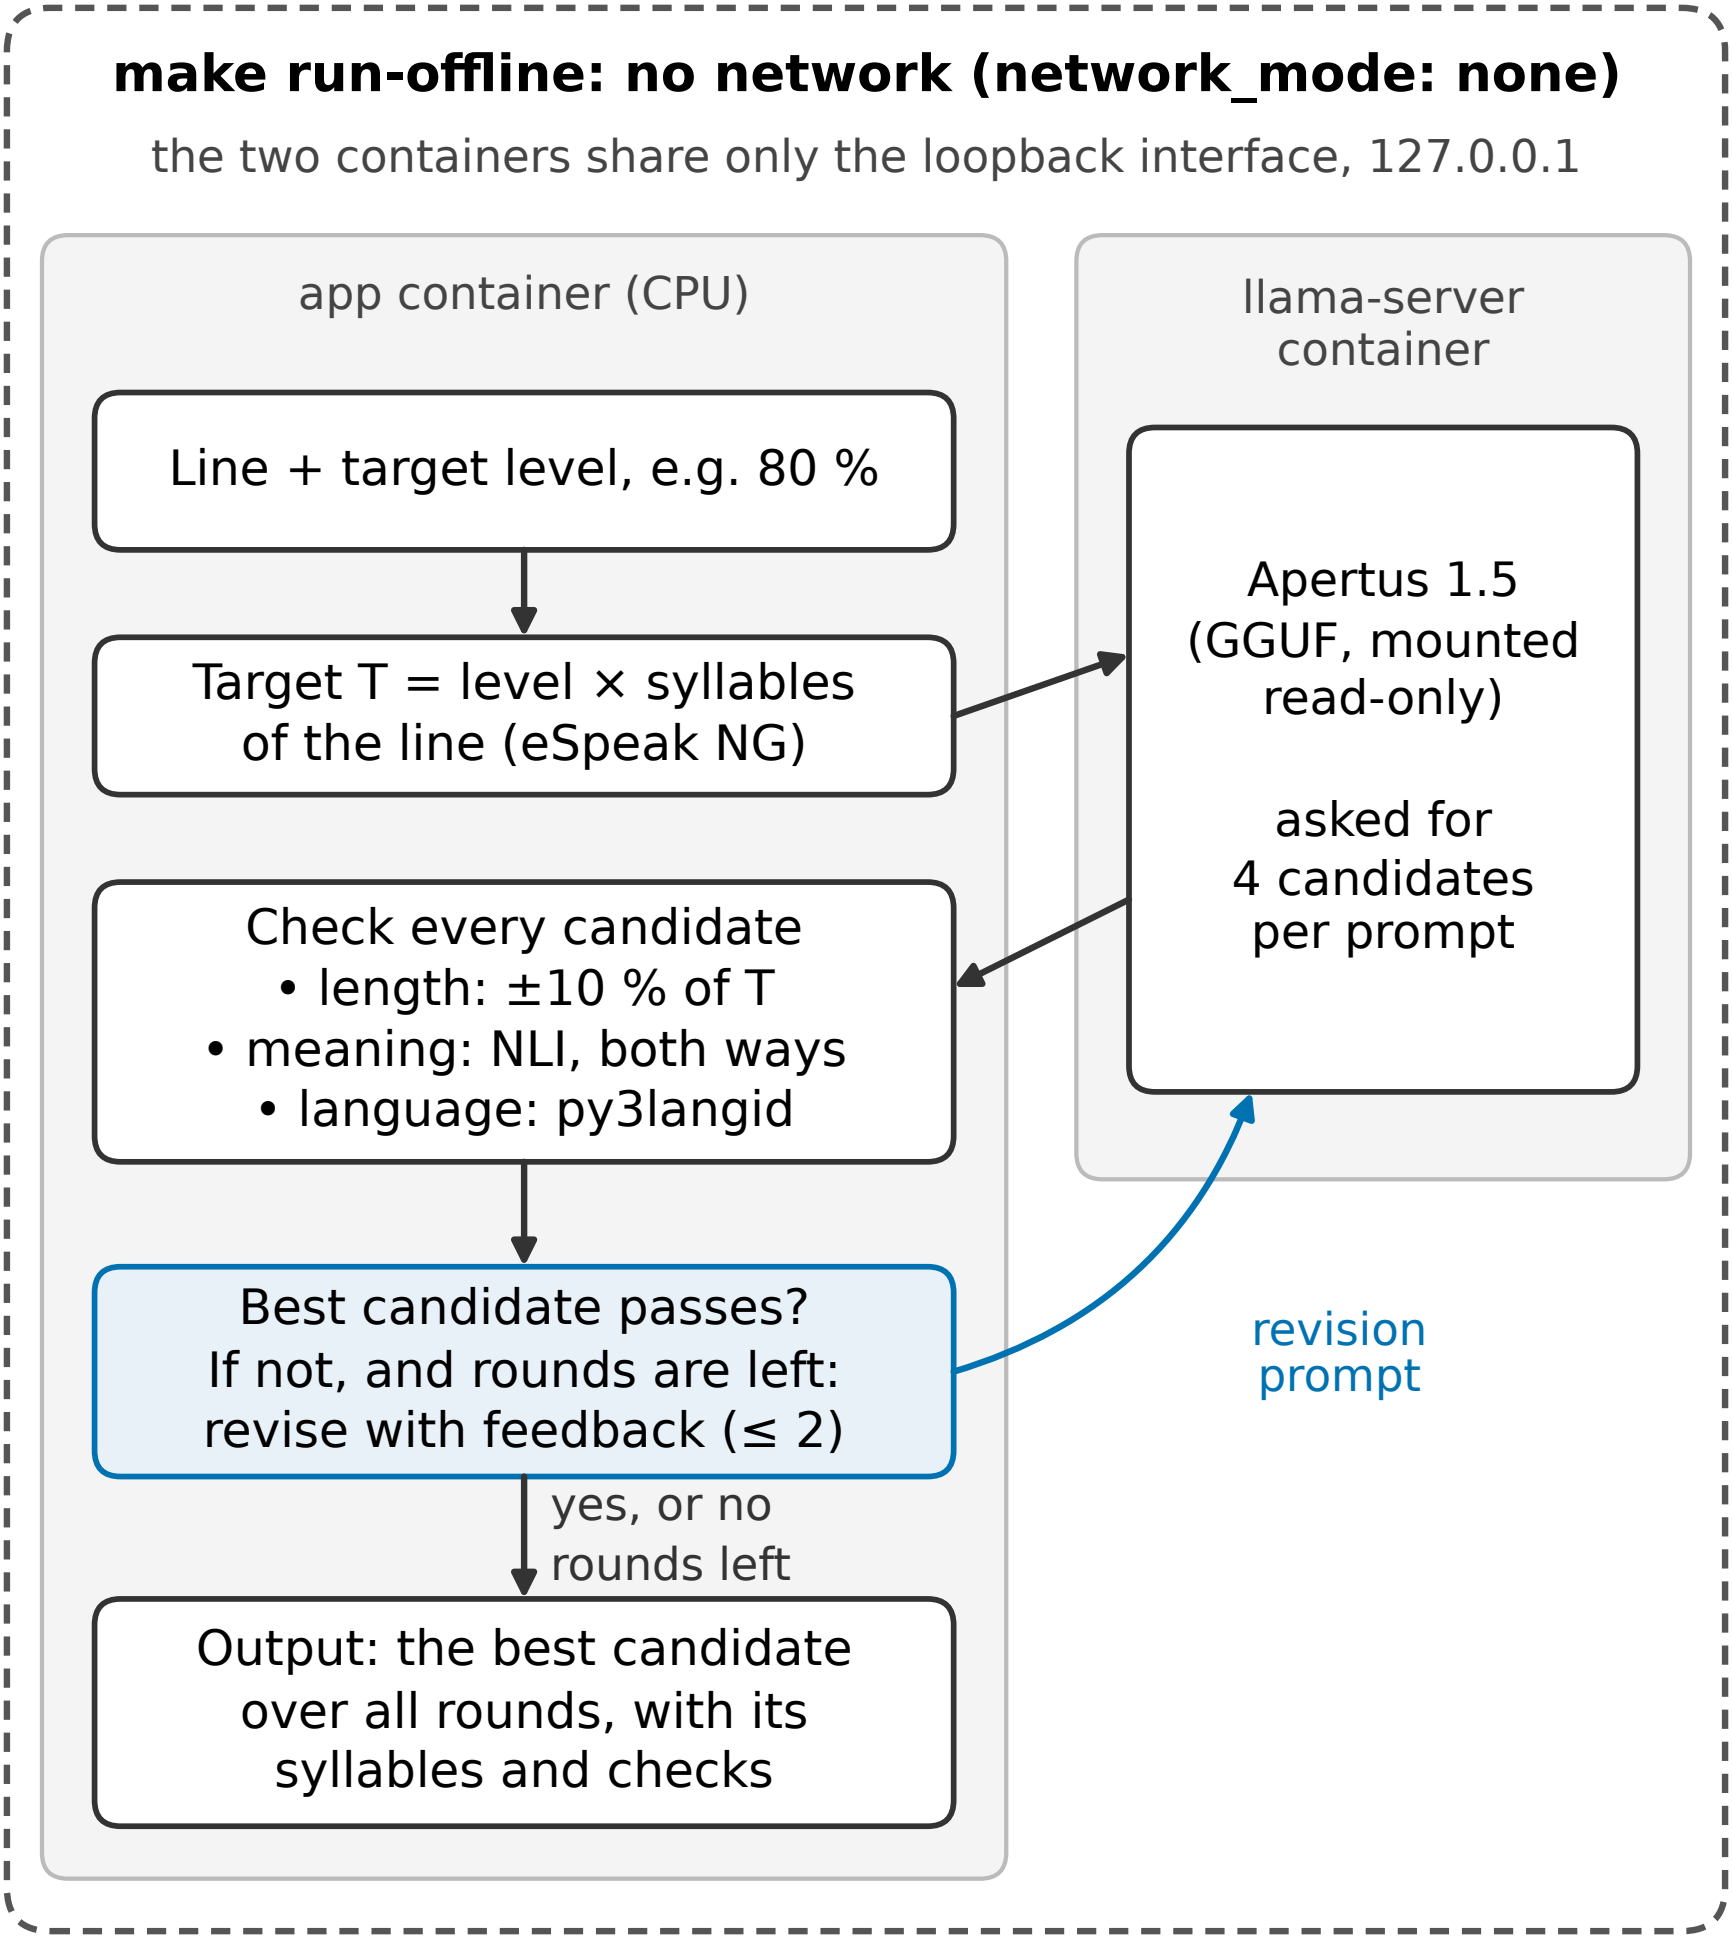
make run-offline: no network (network_mode: none)
the two containers share only the loopback interface, 127.0.0.1
app container (CPU)
llama-server container
Line + target level, e.g. 80 %
Target T = level × syllables of the line (eSpeak NG)
Check every candidate
• length: ±10 % of T
• meaning: NLI, both ways
• language: py3langid
Best candidate passes?
If not, and rounds are left:
revise with feedback (≤ 2)
yes, or no rounds left
Output: the best candidate over all rounds, with its syllables and checks
Apertus 1.5 (GGUF, mounted read-only)
asked for 4 candidates per prompt
revision prompt